

k-Nearest Neighbors (k-NN) Implementation in Parallel

Epameinondas Bakoulas and Maria Sotiria Kostomanolaki

November 2024

Abstract

This document explains the implementation of the **k-Nearest Neighbors (kNN)** algorithm in C, both in serial and parallel execution.

The primary focus is on parallel execution using threads to accelerate performance, particularly for large datasets. We consider scenarios where $C = Q$, the number of data points reaches millions and the dimensionality is in the hundreds.

Additionally, the document explores efficient memory management techniques, such as data block partitioning, to optimize resource utilization.

1 Distance calculation

In a d -dimensional space, the **Euclidean distance matrix** D between two datasets C and Q is calculated as:

$$D = \sqrt{C.^2 - 2CQ^T + (Q.^2)^T}$$

where:

- C is the **corpus** dataset with m points and d dimensions
- Q is the **query** dataset with n points and d dimensions

and $C.^2$ and $Q.^2$ denote element-wise squaring of the matrices and summing along the columns.

The function that implements the above is called **computeDistances** and it is called to calculate the distances between the points in the query and corpus datasets, which initialize the D array of size $m \times n$.

2 Sequential implementation

2.1 Data structures

- **Neighbor** struct: Represents a single neighbor of a data point. Contains an **index** and a **distance**.
- **nearestNeighbors** array: An $n \times k$ array of **Neighbor** structures that stores the k -nearest neighbors for each point in the query dataset. During initialization, the values of **distance** and **index** are set to *INFINITY* and -1 respectively.
- **D** array: An $m \times n$ array that stores the distances between each pair of points in the query and corpus datasets.

2.2 Algorithm

The **kNNsearch** function finds the k -nearest neighbors in matrix C for each point in matrix Q .

To optimize memory usage, Q is divided into smaller *blocks* ($Q1, Q2, \dots$). The distances between query points in *each block* and *all* points in C are then computed using the **computeDistances** function. Finally, the **quickSelect** function is called to efficiently determine the k -nearest neighbors for the current query and their corresponding distances. The results are stored in the **nearestNeighbors** array, after they are compared against previously stored neighbors.

To match the behavior of MATLAB's **knnsearch** function, the **quickSelect** function is modified to sort the neighbors by distance.

Additionally, the output of the **kNNsearch** function is given as a pair of arrays:

- **idx** : Contains the indices of the k-nearest neighbors for each point in Q .
- **dist** : Contains the distances to the corresponding nearest neighbors.

By following this approach, our implementation aims to replicate the functionality and output format of MATLAB's `knnsearch` function.

2.3 Analysis

The sequential approach, while guaranteeing excellent accuracy, can be computationally expensive, particularly for large datasets. This is due to the exhaustive calculation of distances between every pair of points, regardless of their proximity. This indiscriminate computation often involves processing "far" points that ultimately contribute little to the final nearest neighbor determination.

3 Parallel implementation

To efficiently handle large-scale datasets, we adopt a parallel computing approach that prioritizes speed over absolute precision. By employing *approximation* techniques, we aim to achieve a balance between computational efficiency and solution quality.

To simplify our analysis, we will assume that $C == Q$. We will consider datasets with millions of points and dimensions between 3 and 1000.

The implementation of the algorithm is done using the `Pthreads`, `OpenMP`, and `OpenCilk` libraries.

3.1 Data structures

The data structures used in the parallel implementation are similar to the sequential version, with the addition of a few key elements:

- **shuffledIndices** array: Contains the shuffled indices of the points in the dataset.

The initializations of the `nearestNeighbors` and `D` arrays are the same as in the sequential implementation.

3.2 Splitting data into blocks

Controlling memory usage is crucial, so we split the data randomly into *blocks*. The randomness is achieved with the helper function `shuffleIndices` that takes an array of indices and shuffles them randomly. The first $n/\text{numBlocks}$ shuffled points are assigned to the first block, the next $n/\text{numBlocks}$ points to the second block, and so on.

For example, consider the matrix C ($n \times d$):

$$C = \begin{bmatrix} 0 & 1 \\ 1 & 2 \\ 2 & 3 \\ 3 & 4 \\ 4 & 5 \\ 5 & 6 \\ 6 & 7 \\ 7 & 8 \\ 8 & 9 \\ 9 & 10 \\ 10 & 11 \\ 11 & 12 \\ 12 & 13 \\ 13 & 14 \\ 14 & 15 \\ 15 & 16 \end{bmatrix}$$

one possible split into 4 blocks is:

$$\text{Block 1} = \begin{bmatrix} 0 & 1 \\ 5 & 6 \\ 10 & 11 \\ 15 & 16 \end{bmatrix}, \quad \text{Block 2} = \begin{bmatrix} 2 & 3 \\ 4 & 5 \\ 12 & 13 \\ 14 & 15 \end{bmatrix}, \quad \text{Block 3} = \begin{bmatrix} 1 & 2 \\ 7 & 8 \\ 9 & 10 \\ 13 & 14 \end{bmatrix}, \quad \text{Block 4} = \begin{bmatrix} 3 & 4 \\ 6 & 7 \\ 8 & 9 \\ 11 & 12 \end{bmatrix}$$

3.3 Processing blocks

The kNN algorithm is implemented in a **block-wise** manner to improve efficiency. Each block of data points is processed independently in parallel.

For each block, pairwise distances between points are calculated using the `computeDistances` function. To efficiently identify the k-Nearest Neighbors for a given point, the `quickSelect` function is called on an array of `Neighbor` structures, each containing a `distance` and an `index`. The k-Nearest Neighbors are then updated in the `nearestNeighbors` array.

After processing all of the blocks, the `nearestNeighbors` matrix (NN) will look like this (only first 4 points are shown, $k = 3$):

$$\text{NN} = \begin{bmatrix} (0, 0.0) & (5, 7.07) & (10, 14.14) \\ (1, 0.0) & (7, 8.49) & (9, 11.31) \\ (2, 0.0) & (4, 2.83) & (12, 14.14) \\ (3, 0.0) & (6, 4.24) & (8, 7.07) \end{bmatrix}$$

3.4 Improving the Solution

The solution is improved by finding distances between points in different blocks using a *subset* of the points. For example, if we combine pairs of blocks by choosing only 50% of the points in random (2 points from each block), we might get:

$$\text{Subset from Block 1} = \begin{bmatrix} 0 & 1 \\ 10 & 11 \end{bmatrix}, \quad \text{Subset from Block 2} = \begin{bmatrix} 2 & 3 \\ 14 & 15 \end{bmatrix}$$

We will calculate the distance matrix using the points $[0, 1]$, $[2, 3]$, $[10, 11]$ and $[14, 15]$ to update their neighbors. For each point, we will find the distances to the points in the opposite sub-block. This can be achieved using the `computeDistances` function by passing the first subset as the **corpus** and the second subset as the **query**. This approach ensures that we do not calculate the distances between points within the same sub-block.

Updating the nearest neighbors is achieved with the function `updateKNearestNeighbors`. It finds the (current) nearest neighbor with the biggest distance, and tries to replace it with the new neighbor. If the new neighbor is closer, it replaces the current one.

After this improvement, the `nearestNeighbors` matrix (NN) will look like this:

$$\text{NN} = \begin{bmatrix} (0, 0.0) & (5, 7.07) & (2, 2.83) \\ (1, 0.0) & (7, 8.49) & (9, 11.31) \\ (2, 0.0) & (4, 2.83) & (0, 2.83) \\ (3, 0.0) & (6, 4.24) & (8, 7.07) \end{bmatrix}$$

The block pairs are processed in parallel. Each thread processes a pair of blocks, until all pairs are processed. If we process the pair (i, j) , we do not have to process the pair (j, i) , since the distances are symmetric.

3.5 Analyzing the end results

The last step is to find the *Recall* and *Queries per second* of the algorithm. Recall is equal to the percentage of the correct neighbors that we find. It is calculated by comparing the ground truth (exact) solution of the algorithm with the approximate solution.

$$\text{Recall} = \frac{\# \text{ correct neighbors}}{\text{total neighbors}}$$

Queries per second is the number of queries that the algorithm processes in one second.

$$\text{Queries per second} = \frac{\# \text{ points}}{\text{execution time (seconds)}}$$

4 Benchmarks

We will display the benchmarks produced from testing the dataset *SIFT-128-euclidean*. It has a total of $n = 1,000,000$ points and $d = 128$ dimensions, and we'll find the $k = 100$ nearest neighbors. Tested on a computer with CPU: i5-11400F (6 cores, 12 threads), 16GB RAM, and OS: Linux Mint 22. The following graph summarizes the benchmarks for different implementations:

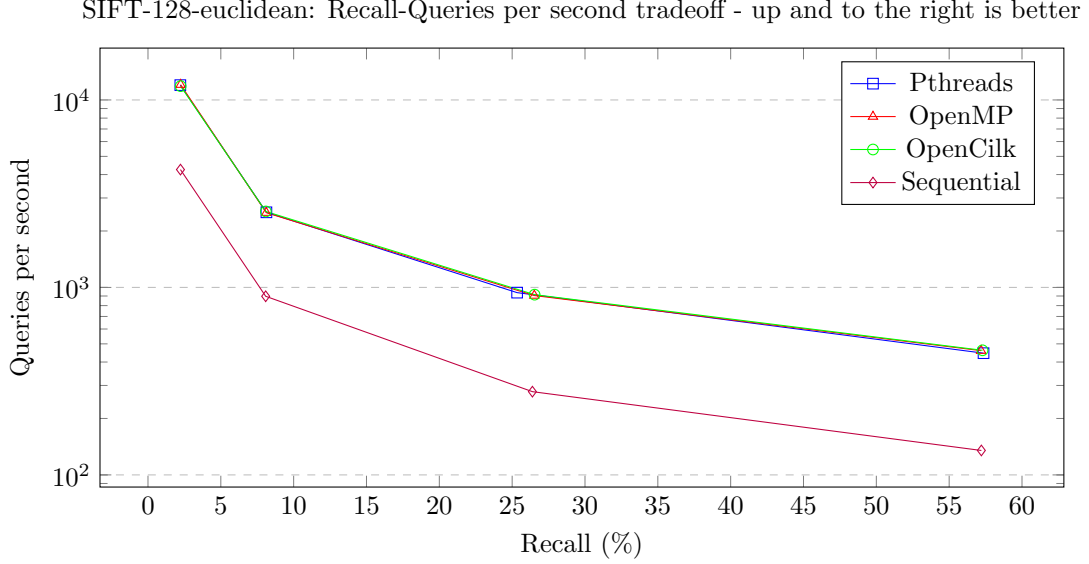


Figure 1: Queries per second vs Recall for different implementations (4 threads, 100 blocks)

All 3 parallel implementations produce very similar results. We can clearly see that the parallel implementations are much faster than the sequential one by a factor of **x2.8 - 3.4** using 4 threads.

To measure the recall we used only the first 10000 points of the dataset and we did not take into account the position of the neighbors in the `nearestNeighbors` array. We only checked if the neighbors were the same with the results produced by Matlab's `knnsearch` function. The correctness of this approach is guaranteed by the fact that we introduced randomness in the algorithm, and also because n is very large.

What if we wanted to maximize the *Recall* for this dataset using the same algorithm? We tested it with OpenMP and the end results were Recall = 99.98% and Queries per second = 241. This result beats the last sequential result (Recall = 57.21%, Queries per second = 135) by a factor of $\times 1.7$, both in Recall and Queries per second.

We can achieve even better performance by increasing the number of threads. Testing the same dataset with 8 threads the results were slightly better, by a factor of $\times 1.1 - 1.2$ in Queries per second, compared to 4 threads. This is expected due to hardware bottlenecks, but it does allow us to increase the utilization of the CPU and boost performance.

5 Algorithm correctness

The algorithm does produce correct results without bias, since there's randomness involved throughout the process. Block splitting is completely random, and picking points to create sub blocks is also random. In the case of large datasets, rerunning the algorithm with the same parameters will produce very similar *Recall* results with a very small deviation.

Data races may occur since `nearestNeighbors` is a global variable shared between the threads, but the error they might produce is marginal, compared to the overall result. The algorithm was tested with locks implemented to avoid data races, but it came at a high cost of speed. Thus, in this case, the speed of the algorithm takes a higher precedence.

6 Summary

The parallel algorithm can be briefly summarized as:

1. Split the data points in blocks (ex. 100 blocks, with 10k points each if we have 1M points).
2. Within each block, calculate the distances between the points and find the k-NN.
3. Find new neighbors by taking a subset of the points from 2 blocks (ex. 50%), calculate the new distances and update the k-NN if necessary.
4. Repeat step 3 for all block pairs.

Steps 1 and 3 do not split the points sequentially. Instead, they shuffle the indices of the points randomly.

7 Source code

You can find the source code of the implementation on the Github repository.

Detailed instructions on how to run the code and the results of the benchmarks can be found in the repository's README file.